


Per-label CER vs. epoch count (one panel per model, own y-scale)

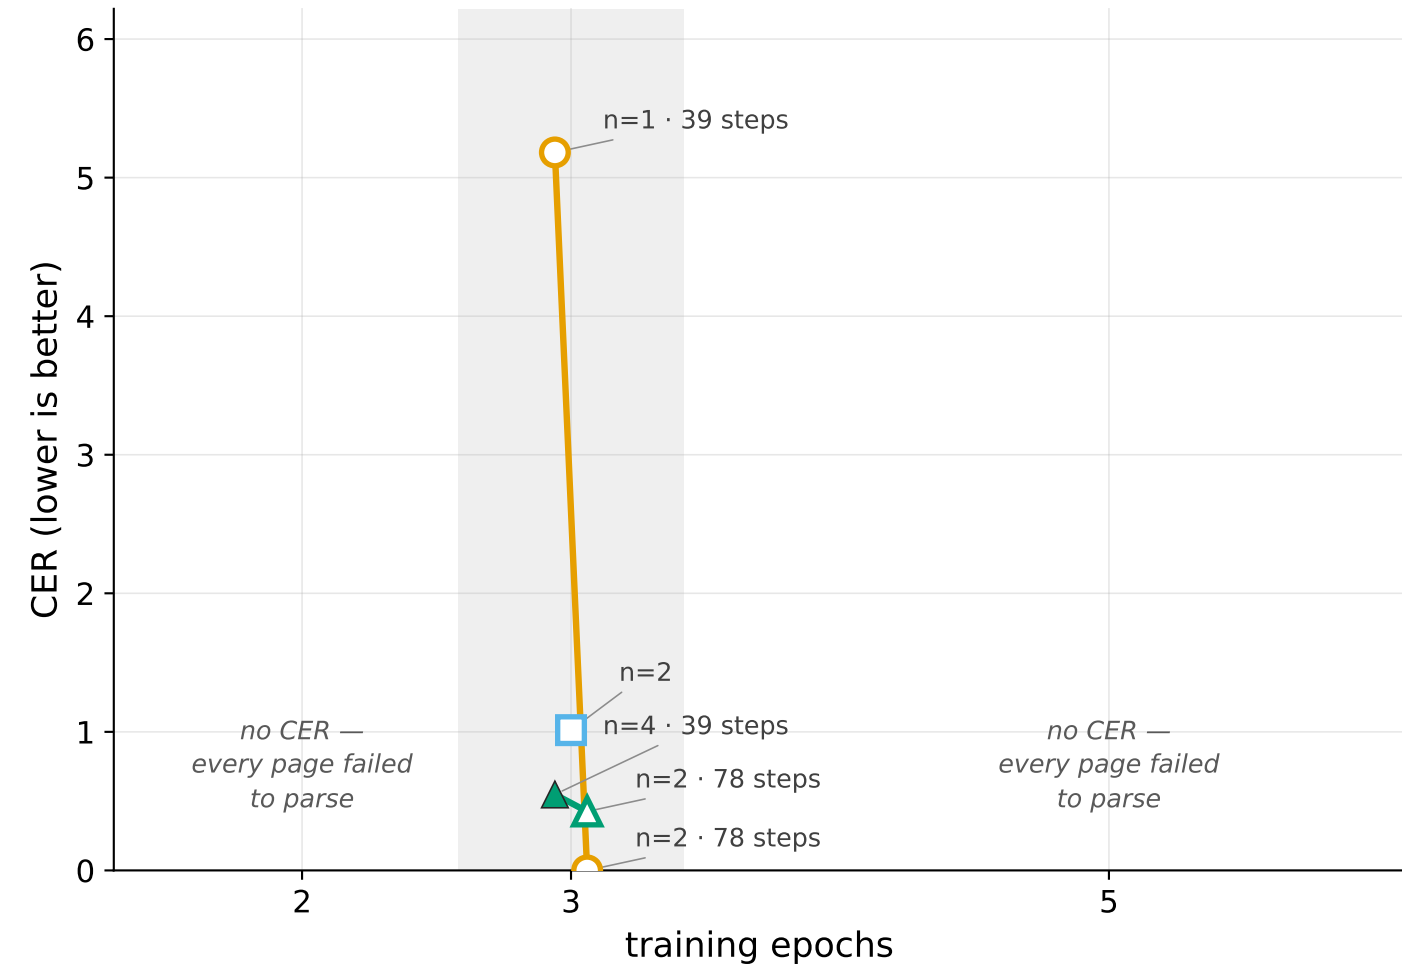
n = pages behind each average. It shrinks as parse failures rise, so the worst-performing runs are also the ones resting on the least data.

Page-Furniture Section-Header Table Text

Gemma 4 E2B



Qwen3.5-0.8B



Note: line style = dataset version (solid = v5, the controlled sweep; dashed = v2, an earlier/smaller dataset, faded); hollow marker = averaged over 3 pages or fewer; Qwen3.5-0.8B's v5, 3-epoch runs differ in step count (39 vs. 78 steps = different effective batch size), not duplicate configs.