



Fine-tuning evaluation — full-detail reference sheet

Every chart here also exists as its own slide-ready file in this folder, drawn from the same code with the per-point labels, provenance and footnotes below stripped out.
Those are the versions to present or embed; this sheet is the backup for a specific follow-up question. Numbers come straight from eval/*.csv — nothing is recomputed here.

1 · Real-document evaluation: same 15 hand-verified ARDB pages, all three approaches

The existing OCR pipeline wins on every metric measured — with no fine-tuning at all.


OCR pipeline (Surya, no fine-tune) Gemma 4 E2B (fine-tuned, v5-e3) Qwen3.5-0.8B (fine-tuned, v5-e3-ga2)



Qwen3.5-0.8B: all 15 pages failed to parse, so it has no bar on either panel (an empty slot is not a score of zero)

2 · JSON parse-failure rate on the validation split vs. epoch count

Both models fail least at 3 epochs on v5; 2 and 5 epochs are worse. One run per point — a direction, not a tested effect.





Line style = dataset version (solid = v5, the controlled sweep; dashed = v2, an earlier/smaller dataset, shown faded for context only). Each point is one run scored on 9 validation pages.

3 · Per-label CER on the validation split vs. epoch count (one panel per model, own y-scale)

n = pages behind each average. It shrinks as parse failures rise, so the worst-performing runs are also the ones resting on the least data.

Page-Furniture Section-Header Table Text



Note: line style = dataset version (solid = v5, the controlled sweep; dashed = v2, an earlier/smaller dataset, faded); hollow marker = averaged over 3 pages or fewer; Qwen3.5-0.8B's v5, 3-epoch runs differ in step count (39 vs. 78 steps = different effective batch size), not duplicate configs.

4 · Training loss by step, one panel per model

Every run converged cleanly — including the ones whose output was unusable. This is a training-process check, not a quality signal.

