



Fine-tuning experiment at a glance

Two models, three epoch settings each. Both peaked at 3 epochs — and the best of them still lost to the pipeline already in production.

1. The sweep: does the output even parse?



2. The verdict: best adapter vs. the existing pipeline



Left: in-distribution synthetic validation split (9 pages/run, one run per point). Right: 15 held-out, hand-verified real pages.
Single seed throughout — no error bars. Per-run detail, sample sizes and caveats: the four charts that follow, and this folder's README.