



Real-document evaluation: same 15 hand-verified ARDB pages, all three approaches

The existing OCR pipeline wins on every metric measured — with no fine-tuning at all.



Higher is better

Lower is better



Qwen3.5-0.8B: all 15 pages failed to parse, so it has no bar on either panel (an empty slot is not a score of zero)