


Training loss by step, one panel per model

Every run converged cleanly — including the ones whose output was unusable. This is a training-process check, not a quality signal.

Gemma 4 E2B



Qwen3.5-0.8B

